

UMBC CSEE

Foundations of Data Mining

CS-691

Homework Assignment - 4

Sandipan Dey



2009

1. Construct a feed-forward 3-layer (one input layer, one hidden layer, and one output node) perceptron network that correctly learns the XOR function. Clearly show the link weights and the thresholds. Construct this network through analytical observations, not by running a neural network simulation environment.

Answer:

Let's denote the input binary variables by $x_1, x_2 \in \{0,1\}$. We have to learn the binary function $XOR: \{0,1\} \times \{0,1\} \rightarrow \{0,1\}$, so that the output binary variable $y = x_1 \oplus x_2$.

Since we know that XOR is linearly inseparable function, we can't design the perceptron using just 2-layer perceptron, for if we could, $\exists w_1, w_2$ (weights) and a non-negative threshold $t \geq 0$ such that

$$y = \begin{cases} 0, & w_1 x_1 + w_2 x_2 < t \\ 1, & w_1 x_1 + w_2 x_2 \geq t \end{cases}$$

but since $y = x_1 \oplus x_2$, with XOR function defined by the following:



| x_1 | x_2 | y |
|-------|-------|-----|
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |

with all the following inequalities must be satisfied:

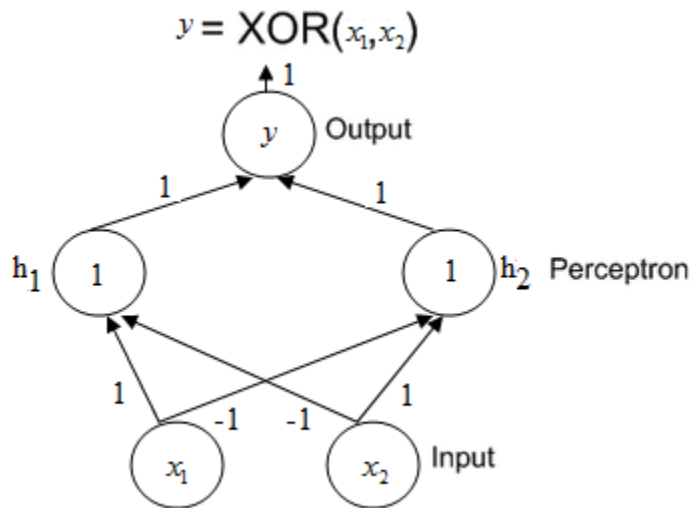
$$w_1 \geq t, \quad w_2 \geq t, \quad 2t \leq w_1 + w_2 < t, \quad t \geq 0 \Rightarrow t > 2t \wedge t \geq 0 \Rightarrow t < 0 \wedge t \geq 0,$$

a contradiction.

Hence, we need a 3-layer perceptron network, with a hidden layer placed in between the input and output layer.



Now the XOR function can be realized as a linearly separable function of couple of linearly separable functions (in 2 steps): $y = x_1 \oplus x_2 = x_1 \overline{x_2} + \overline{x_1} x_2 = \text{AND}(\text{OR}(x_1, \text{NOT}(x_2)), \text{OR}(\text{NOT}(x_1), x_2))$, all of AND, OR, NOT functions being linearly separable.




Threshold is 1 everywhere at both the nodes in the hidden layer and also at the output node.

| x_1 | x_2 | h_1 | h_2 | y |
|-------|-------|-------|-------|-----|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 0 | 1 | 1 |
| 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 0 | 0 | 0 |

Alternative design

In order design the following 3-layer perceptron, we see that the problem is due to the (1,1) input, we must make it less than the threshold somehow, to have output 0. We can see that it can be done by associating a negative weight to a node in the hidden layer, when both x_1, x_2 have value 1.



$$\begin{aligned}
2. \quad \phi(x) &= [1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1 x_2, x_1^2, x_2^2] \\
\Rightarrow \phi(u).\phi(v) &= [1, \sqrt{2}u_1, \sqrt{2}u_2, \sqrt{2}u_1 u_2, u_1^2, u_2^2]. [1, \sqrt{2}v_1, \sqrt{2}v_2, \sqrt{2}v_1 v_2, v_1^2, v_2^2]^T \\
\Rightarrow K(u, v) &= \phi(u).\phi(v) = 1 + 2u_1 v_1 + 2u_2 v_2 + 2u_1 v_1 u_2 v_2 + u_1^2 v_1^2 + u_2^2 v_2^2 \\
\Rightarrow K(u, v) &= \left(\begin{bmatrix} u_1 & u_2 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \end{bmatrix} + 1 \right)^2 = (u.v + 1)^2
\end{aligned}$$

Now, the dual problem becomes

$$\begin{aligned}
L_D(\lambda) &= \sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j \phi(x_i) . \phi(x_j) \\
\Rightarrow L_D(\lambda) &= \sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j K(x_i, x_j)
\end{aligned}$$

We have the following training points:

$$x_1=(1, 1, -), x_2=(1, 0, +), x_3=(0, 1, +), x_4=(0, 0, -), y_1=y_3=+1, y_2=y_4=-1.$$

Assuming + label to be with value +1 and the - label with value -1, let's first construct the Kernel matrix:

$$K(x_1, x_1) = ([1 \ 1]. [1 \ 1]^T + 1)^2 = (2 + 1)^2 = 9,$$

$$K(x_1, x_2) = ([1 \ 1]. [1 \ 0]^T + 1)^2 = (1 + 1)^2 = 4$$

$$K(x_2, x_3) = ([1 \ 0]. [0 \ 1]^T + 1)^2 = (0 + 1)^2 = 1$$

| | Kernel | (1,1) | (1,0) | (0,1) | (0,0) |
|-------|--------|-------|-------|-------|-------|
| x_1 | (1, 1) | 9 | 4 | 4 | 1 |
| x_2 | (1, 0) | 4 | 9 | 1 | 1 |
| x_3 | (0, 1) | 4 | 1 | 9 | 1 |
| x_4 | (0, 0) | 1 | 1 | 1 | 9 |

$$\sum_{i=1}^4 \sum_{j=1}^4 \lambda_i \lambda_j y_i y_j K(x_i, x_j) = 9\lambda_1^2 - 4\lambda_1 \lambda_2 + 4\lambda_1 \lambda_3 - \lambda_1 \lambda_4$$

$$\begin{aligned}
&- 4\lambda_1 \lambda_2 + 9\lambda_2^2 - \lambda_2 \lambda_3 + \lambda_2 \lambda_4 + 4\lambda_1 \lambda_3 - \lambda_2 \lambda_3 + 9\lambda_3^2 - \lambda_3 \lambda_4 - \lambda_1 \lambda_4 + \lambda_2 \lambda_4 - \lambda_3 \lambda_4 + 9\lambda_4^2 \\
&= 9(\lambda_1^2 + \lambda_2^2 + \lambda_3^2 + \lambda_4^2) - 8\lambda_1 \lambda_2 + 8\lambda_1 \lambda_3 - 2\lambda_1 \lambda_4 - 2\lambda_2 \lambda_3 + 2\lambda_2 \lambda_4 - 2\lambda_3 \lambda_4
\end{aligned}$$

$$L_D = \sum_{i=1}^4 \lambda_i - \frac{1}{2} \sum_{i=1}^4 \sum_{j=1}^4 \lambda_i \lambda_j y_i y_j K(x_i, x_j)$$

$$= (\lambda_1 + \lambda_2 + \lambda_3 + \lambda_4) - (9/2)(\lambda_1^2 + \lambda_2^2 + \lambda_3^2 + \lambda_4^2) + 4\lambda_1 \lambda_2 - 4\lambda_1 \lambda_3 + \lambda_1 \lambda_4 + \lambda_2 \lambda_3 - \lambda_2 \lambda_4 + \lambda_3 \lambda_4$$

Taking partial derivative w.r.t . λ_i for optimizing L_D , we get the following 4 equations:

$$9\lambda_1 - 4\lambda_2 + 4\lambda_3 - \lambda_4 = 1$$

$$-4\lambda_1 + 9\lambda_2 - \lambda_3 + \lambda_4 = 1$$

$$4\lambda_1 - \lambda_2 + 9\lambda_3 - \lambda_4 = 1$$

$$-\lambda_1 + \lambda_2 - \lambda_3 + 9\lambda_4 = 1$$

With solution: $\lambda_1 = 0.1765, \lambda_2 = 0.1838, \lambda_3 = 0.0662, \lambda_4 = 0.1176$

Also, $\phi(x) = [1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2, x_1^2, x_2^2]$

$$\Rightarrow \phi(x_1) = [1, \sqrt{2}, \sqrt{2}, \sqrt{2}, 1, 1]$$

$$\Rightarrow \phi(x_2) = [1, \sqrt{2}, 0, 0, 1, 0]$$

$$\Rightarrow \phi(x_3) = [1, 0, \sqrt{2}, 0, 0, 1]$$

$$\Rightarrow \phi(x_4) = [1, 0, 0, 0, 0, 0]$$

$$\Rightarrow w_o = \sum_i \lambda_i y_i \phi(x_i) = [-0.5690 \quad -0.3415 \quad -0.5078 \quad -0.6014 \quad 0.0063 \quad 0.6342]$$

$$\Rightarrow \max margin = (1/2) \|w_o\|^2 = 0.5 \times 1.2091 = 0.6046 \text{ unit}$$

3. (a) Applying Bayes theorem,

$$P(A=1|+) = \frac{P(+|A=1)P(A=1)}{P(+)} = \frac{(3/5).(1/2)}{(1/2)} = \frac{3}{5}$$

$$P(B=1|+) = \frac{P(+|B=1)P(B=1)}{P(+)} = \frac{(2/4).(4/10)}{(1/2)} = \frac{2}{5}$$

$$P(C=1|+) = \frac{P(+|C=1)P(C=1)}{P(+)} = \frac{(4/5).(1/2)}{(1/2)} = \frac{4}{5}$$

$$P(A=1|-) = \frac{P(-|A=1)P(A=1)}{P(-)} = \frac{(2/5).(1/2)}{(1/2)} = \frac{2}{5}$$

$$P(B=1|-) = \frac{P(-|B=1)P(B=1)}{P(-)} = \frac{(2/4).(4/10)}{(1/2)} = \frac{2}{5}$$

$$P(C=1|-) = \frac{P(-|C=1)P(C=1)}{P(-)} = \frac{(1/5).(1/2)}{(1/2)} = \frac{1}{5}$$

(b) $P(+|A=1, B=1, C=1)$

$$\begin{aligned} &= \frac{P(A=1, B=1, C=1|+)P(+)}{P(A=1, B=1, C=1)} = \frac{P(A=1|+)P(B=1|+)P(C=1|+)P(+)}{P(A=1, B=1, C=1)} \\ &= \frac{(3/5).(2/5).(4/5).(1/2)}{1/10} = \frac{3.1.4.10}{5.5.5} = \frac{24}{25} \end{aligned}$$

$P(-|A=1, B=1, C=1)$

$$\begin{aligned} &= \frac{P(A=1, B=1, C=1|-)P(-)}{P(A=1, B=1, C=1)} = \frac{P(A=1|-)P(B=1|-)P(C=1|-)P(-)}{P(A=1, B=1, C=1)} \\ &= \frac{(2/5).(2/5).(1/5).(1/2)}{1/10} = \frac{1.2.1.10}{5.5.5} = \frac{4}{25} \end{aligned}$$

Hence, prediction using Naive Bayesian is class label +.

(c) $P(A = 1) = 1/2$.

$P(B = 1) = 4/10 = 2/5$.

$P(A = 1, B=1) = 2/10 = 1/5$.

Since $P(A=1, B=1) = 1/5 = (1/2) \cdot (2/5) = P(A=1)P(B=1)$, A and B are statistically independent.

4. The popular approaches to handle missing values are the following:

a) *Discard instances with missing values*

b) *Imputation (by mean substitution, multiple regression, MLE, nearest neighbor etc)*

Here we are going to apply both the techniques discard /mean substitution separately on the data.

We use 10 fold cross-validation technique. To measure performance of the model, we use the following metrics:

■ Accuracy = $(TP+TN)/(TP+TN+FP+FN)$
■ Precision = $TP/(TP+FP)$
■ Recall = $TP/(TP+FN)$

Where TP = True Positive, FP = False Positive

TN = True Negative, FN = False Negative

Using **Naïve Bayesian Classifier**, we obtain the following results using weka:

Time taken to build model: 0.01 seconds

=== Stratified cross-validation ===

| | | |
|----------------------------------|-----------|---|
| Correctly Classified Instances | 515 | 78.0303 % |
| Incorrectly Classified Instances | 145 | 21.9697 % |
| Kappa statistic | 0.5439 | |
| Mean absolute error | 0.2212 | |
| Root mean squared error | 0.434 | |
| Relative absolute error | 44.6569 % | |
| Root relative squared error | 87.2083 % | |
| Total Number of Instances | 660 | (after discarding tuples with missing values) |

=== Detailed Accuracy By Class ===

| | TP Rate | FP Rate | Precision | Recall | F-Measure | ROC Area | Class |
|---------------|---------|---------|-----------|--------|-----------|----------|-------|
| | 0.601 | 0.072 | 0.873 | 0.601 | 0.712 | 0.901 | + |
| | 0.928 | 0.399 | 0.738 | 0.928 | 0.823 | 0.902 | - |
| Weighted Avg. | 0.78 | 0.251 | 0.799 | 0.78 | 0.772 | 0.902 | |

=== Confusion Matrix ===

a b <-- classified as

179 119 | a = +

26 336 | b = -

with TP = 179, TN = 336, FP = 26, FN = 119.

$$\text{Hence, Accuracy} = \frac{TP + TN}{TP + FP + TN + FN} = \frac{179 + 336}{179 + 26 + 336 + 119} = 0.78$$

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{179}{179 + 26} = 0.83$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{179}{179 + 119} = 0.60$$

=== Stratified cross-validation ===

| | | |
|----------------------------------|-----------|---|
| Correctly Classified Instances | 537 | 77.8261 % |
| Incorrectly Classified Instances | 153 | 22.1739 % |
| Kappa statistic | 0.5372 | |
| Mean absolute error | 0.2227 | |
| Root mean squared error | 0.4355 | |
| Relative absolute error | 45.0874 % | |
| Root relative squared error | 87.634 % | |
| Total Number of Instances | 690 | (without discarding tuples with missing values) |

=== Detailed Accuracy By Class ===

| | TP Rate | FP Rate | Precision | Recall | F-Measure | ROC Area | Class |
|---------------|---------|---------|-----------|--------|-----------|----------|-------|
| | 0.599 | 0.078 | 0.86 | 0.599 | 0.706 | 0.896 | + |
| | 0.922 | 0.401 | 0.742 | 0.922 | 0.822 | 0.896 | - |
| Weighted Avg. | 0.778 | 0.257 | 0.794 | 0.778 | 0.77 | 0.896 | |

=== Confusion Matrix ===

a b <-- classified as

184 123 | a = +

30 353 | b = -

with TP = 184, TN = 353, FP = 30, FN = 123.

$$\text{Hence, Accuracy} = \frac{TP + TN}{TP + FP + TN + FN} = \frac{183 + 353}{183 + 30 + 353 + 123} = 0.78$$

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{184}{184 + 30} = 0.86$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{184}{184 + 123} = 0.61$$

As we can see, the performance is improved in this case (if we don't discard tuples with missing values).

Using J48 classification algorithm (decision tree)

J48 pruned tree

```
A9 = t
|  A10 = t: + (228.0/21.0)
|  A10 = f
|  |  A15 <= 444
|  |  |  A7 = v
|  |  |  |  A4 = u
|  |  |  |  |  A14 <= 112: + (16.57/1.57)
|  |  |  |  |  A14 > 112
|  |  |  |  |  |  A15 <= 70: - (30.0/10.0)
|  |  |  |  |  |  A15 > 70: + (2.0)
|  |  |  |  |  A4 = y
|  |  |  |  |  |  A13 = g: - (12.0/2.0)
|  |  |  |  |  |  A13 = s: + (3.0/1.0)
|  |  |  |  |  |  A13 = p: - (0.0)
|  |  |  |  |  A4 = l: - (0.0)
|  |  |  |  |  A7 = h: + (27.24/8.24)
|  |  |  |  |  A7 = bb
|  |  |  |  |  A3 <= 1.375: + (5.0/1.0)
|  |  |  |  |  A3 > 1.375: - (9.13/1.0)
|  |  |  |  |  A7 = ff: - (5.05/1.0)
|  |  |  |  |  A7 = j: - (1.01)
|  |  |  |  |  A7 = z: + (0.0)
|  |  |  |  |  A7 = o: + (0.0)
|  |  |  |  |  A7 = dd: + (1.01/0.01)
|  |  |  |  |  A7 = n: + (0.0)
|  |  |  |  |  A15 > 444: + (21.0/1.0)
A9 = f
|  A3 <= 0.165
|  |  A7 = v
|  |  |  A2 <= 35.58: - (18.72/3.44)
|  |  |  A2 > 35.58: + (3.6/0.16)
|  |  |  A7 = h: - (0.0)
|  |  |  A7 = bb: + (1.24/0.08)
|  |  |  A7 = ff: - (4.96/0.64)
|  |  |  A7 = j: + (1.24/0.08)
|  |  |  A7 = z: - (0.0)
|  |  |  A7 = o: - (0.0)
|  |  |  A7 = dd: - (0.0)
|  |  |  A7 = n: + (1.24/0.08)
|  |  A3 > 0.165: - (298.0/12.0)
```

Time taken to build model: 0.06 seconds

=== Stratified cross-validation ===

| | | |
|----------------------------------|-----------|-----------|
| Correctly Classified Instances | 595 | 86.2319 % |
| Incorrectly Classified Instances | 95 | 13.7681 % |
| Kappa statistic | 0.7208 | |
| Mean absolute error | 0.1907 | |
| Root mean squared error | 0.3303 | |
| Relative absolute error | 38.6088 % | |
| Root relative squared error | 66.4545 % | |
| Total Number of Instances | 690 | |

=== Detailed Accuracy By Class ===

| | TP Rate | FP Rate | Precision | Recall | F-Measure | ROC Area | Class |
|---------------|---------|---------|-----------|--------|-----------|----------|-------|
| | 0.837 | 0.117 | 0.851 | 0.837 | 0.844 | 0.886 | + |
| | 0.883 | 0.163 | 0.871 | 0.883 | 0.877 | 0.886 | - |
| Weighted Avg. | 0.862 | 0.143 | 0.862 | 0.862 | 0.862 | 0.886 | |

=== Confusion Matrix ===

```

a  b  <-- classified as
257 50 | a = +
45 338 | b = -

```

$$\text{Hence, Accuracy} = \frac{TP + TN}{TP + FP + TN + FN} = \frac{257 + 338}{257 + 50 + 338 + 45} = 0.86$$

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{257}{257 + 50} = 0.83$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{257}{257 + 45} = 0.85$$

As we can see, the performance is still improved in this case.

Linear Regression Model

Class =

```

0.1246 * A5=g,gg +
0.8492 * A6=d,i,k,j,aa,m,c,w,e,q,r,cc,x +
-0.6328 * A6=j,aa,m,c,w,e,q,r,cc,x +
0.6058 * A6=aa,m,c,w,e,q,r,cc,x +
0.1715 * A6=m,c,w,e,q,r,cc,x +
-0.3223 * A6=r,cc,x +
0.5857 * A6=cc,x +
-0.6266 * A7=dd,j,v,bb,o,n,h,z +
0.731 * A7=j,v,bb,o,n,h,z +
-0.7769 * A7=v,bb,o,n,h,z +

```

0.5557 * A7=n,h,z +
 -0.4779 * A7=h,z +
 1.1705 * A9=t +
 0.2488 * A10=t +
 0.0153 * A11 +
 1.0218 * A13=p +
 -0.0004 * A14 +
 0 * A15 +
 -1.2233

Time taken to build model: 0.27 seconds

=== Cross-validation ===

| | |
|-----------------------------|-----------|
| Correlation coefficient | 0.7437 |
| Mean absolute error | 0.4588 |
| Root mean squared error | 0.6672 |
| Relative absolute error | 46.329 % |
| Root relative squared error | 66.9622 % |
| Total Number of Instances | 690 |

Multilayer Perceptron

Time taken to build model: 103.25 seconds

=== Cross-validation ===

| | |
|-----------------------------|------------|
| Correlation coefficient | 0.521 |
| Mean absolute error | 0.6858 |
| Root mean squared error | 1.0861 |
| Relative absolute error | 69.2607 % |
| Root relative squared error | 108.9969 % |
| Total Number of Instances | 690 |